

Statistical Learning Theory LAB

Project Report

Maurice Wenig

Contents

1	Algorithms	2
1.1	Neighbourhood	2
1.1.1	User-Based	2
1.1.2	Item-Based	3
1.1.3	Clustering	3
1.2	Factorization	3
1.3	Hybrid	4
2	Results	6
2.1	Performance	7
2.2	Error Scores	7
2.3	Final Score	7

1 Algorithms

1.1 Neighbourhood-Based Recommenders

1.1.1 User-Based Recommender

The user-based recommender compares users based on their rated items. An item rating from a user is then predicted based on ratings of that item from similar users. The similarity measure used to compare users is Pearson:

$$\text{Sim}(u, v) = \frac{\sum_{i \in I_{uv}} w_i \cdot s_{ui} s_{vi}}{\sqrt{\sum_{i \in I_{uv}} w_i \cdot s_{ui}} \cdot \sqrt{\sum_{i \in I_{uv}} w_i \cdot z_{vi}}}$$

where I_{uv} are the items that both users u and v have rated, $s_{ui} = r_{ui} - \mu_u$ is the centered rating from user u . The mean rating μ_u can be calculated individually for every comparison, based on only the common items, or it can be calculated once for every user, based on all their rated items. $w_i = \log \frac{\# \text{ users}}{\# \text{ users who rated } i}$ are item weights to combat the impact of the long tail in the number of item ratings.

To prefer users with more common items, the similarities of users with less than β common items are made smaller.

$$\text{Sim}(u, v) \leftarrow \text{Sim}(u, v) \cdot \frac{\min\{|I_{uv}|, \beta\}}{\beta}$$

The hyper parameter α can be used to amplify the importance of similarity.

$$\text{Sim}(u, v) \leftarrow \text{Sim}(u, v)^\alpha$$

Based on these similarities, a peer group $P_u(i)$ of users is determined. For this, the users are sorted by similarity. To improve efficiency in the online phase, the sorting is done in the offline phase. For this, the similarities between every user u, v are saved in the similarity matrix. The order of peers for a specific user is then determined by sorting the respective row vector.

In the online phase, this order is simply applied. Users are then removed, if they have not rated item i , or their similarity to user u is below a certain threshold S_{min} . The top k remaining users are then the peer group $P_u(i)$.

The predicted item rating from a user is then calculated as a sum of ratings from similar users, weighted with their relative similarity. The means and variances of the user ratings are normalized to not influence the prediction.

$$\hat{r}_{ui} = \mu_u + \sigma_u \cdot \frac{\sum_{v \in P_u(i)} \text{Sim}(u, v) \cdot z_{vi}}{\sum_{v \in P_u(i)} |\text{Sim}(u, v)|}$$

where $z_{ui} = s_{ui}/\sigma_u$ is the normalized rating, σ_u the variance of the ratings of user u . If the user has only ever given one distinct rating, the variance is set to an arbitrary value of 1.

In special cases where the user or the item has never been seen before, similarities can not be calculated. In those special cases, the ratings were predicted in a different way:

- user and item unknown: predict average between minimum and maximum possible rating
- user unknown, item known: predict average rating of that item

- user known, item unknown: predict average rating of that user

Also, if the original item rating of the user is known, this original rating can be predicted.

1.1.2 Item-Based Recommender

The item-based recommender compares items based on their user ratings. An item rating from a user is then predicted based on ratings of similar items from that user. It functions very similar to the user-based recommender. The similarity function is adapted to still use user means μ_u instead of item means μ_i to normalize user preferences. The predicted rating is then a weighted sum of similar items, the user has rated.

$$\hat{r}_{ui} = \frac{\sum_{j \in P_i(u)} \text{Sim}(i, j) \cdot r_{uj}}{\sum_{j \in P_i(u)} |\text{Sim}(i, j)|}$$

1.1.3 Clustering

To increase efficiency, the users/items can be clustered before the creation of the similarity matrix. This way, only users/items that are in the same cluster have to be compared.

This is done using an adapted version of K-means. To account for the sparsity of the rankings matrix, the entries of the mean of a cluster are calculated using only vectors, where the entries are not missing in the respective dimensions. The distance measure is also adapted. The distance between two vectors is calculated only with the common dimensions, where ratings are not missing. The result is then divided by the number of common dimensions to get a distance measure that is independent of the number of common dimensions.

1.2 Factorization-Based Recommender

In the factorization-based recommender, the users and items are assumed to be in a common latent factor space R^f of user preferences q_i and item characteristics p_u . The rating r_{ui} is then modelled as the user's interest in the item's characteristics. User and item biases are also taken into account.

$$\hat{r}_{ui} = q_i^\top p_u + \mu + b_i + b'_u$$

The goal is then to learn the factors q_i, p_u , such that the squared error function on the ratings of the training set is minimal.

$$\min_{p, q, b} \sum_{(u, i) \in \kappa} (r_{ui} - \hat{r}_{ui})^2 + \lambda \left(\|q_i\|^2 + \|p_u\|^2 + b_i^2 + b'_u{}^2 \right)$$

with κ as the set of (u, i) pairs in the training set, λ as the regularization parameter.

To achieve this minimization, alternating least squares is used. q, p are initiated randomly. Then, in each step, one of them is fixed and the other one is solved optimally. This is done until the change of the vectors in q, p is below a threshold ϵ or a maximum number of iterations is reached.

To optimally solve the minimization at each step, ridge regression is used. Because they are

independent, each of the summands are optimized individually.

$$\begin{aligned} \min_{p_u, b'_u} \sum_{i \in \kappa_u} \left((r_{ui} - b_i - \mu) - (q_i^\top p_u + b'_u) \right)^2 + \lambda \left(\|p_u\|^2 + b'^2_u \right) \\ \min_{q_i, b_i} \sum_{u \in \kappa_i} \left((r_{ui} - b'_u - \mu) - (q_i^\top p_u + b_i) \right)^2 + \lambda \left(\|q_i\|^2 + b_i^2 \right) \end{aligned}$$

with κ_u as the set of items that have been rated by u , κ_i as the set of user that have rated i .

The ridge regression boils down to solving the regularized normal equations:

$$\begin{aligned} (Q_{\kappa_u} Q_{\kappa_u}^\top + \lambda \mathbb{1}) p_u &= Q_{\kappa_u} y_{\kappa_u} \\ (P_{\kappa_i} P_{\kappa_i}^\top + \lambda \mathbb{1}) q_i &= P_{\kappa_i} y_{\kappa_i} \end{aligned}$$

with

$$\begin{aligned} Q_{\kappa_u} &= (q_{i_1}, \dots, q_{i_n}) \\ y_{\kappa_u} &= \vec{r}_{\kappa_u} - \vec{b} \\ &= (r_{ui_1}, \dots, r_{ui_n})^\top - (b_{i_1}, \dots, b_{i_n}) \\ i, \dots, i_n &\in \kappa_u \end{aligned}$$

for determining p_u and

$$\begin{aligned} P_{\kappa_i} &= (p_{u_1}, \dots, p_{u_n}) \\ y_{\kappa_i} &= \vec{r}_{\kappa_i} - \vec{b}' \\ &= (r_{u_1 i}, \dots, r_{u_n i})^\top - (b'_{u_1}, \dots, b'_{u_n}) \\ u_1, \dots, u_n &\in \kappa_i \end{aligned}$$

for determining q_i .

The biases can either be fixed or part of the optimization. If the biases are fixed, μ is the overall mean rating in the training set, $b'_u = \mu_u - \mu$ is the mean user rating centered around the overall mean rating, $b_i = \mu_i - \mu$ is the mean item rating centered around the overall mean rating. If the biases are part of the optimization, they are the affine part of the ridge regression at each step. This affine part can be obtained by appending a constant 1 to the training vectors and reading its respective coefficient.

The special cases are handled in the same way as they are in the [User Based Recommender](#).

1.3 Hybrid Recommender

The hybrid combines an array of K recommenders. The predictions are a linear combination of the predictions of those recommenders.

$$\hat{r}_{ui} = \sum_{k=1}^K \alpha_k \hat{r}_{ui}^{(k)} = \vec{\alpha}^\top \vec{r}_{ui}$$

To learn $\vec{\alpha}$, the training set is separated into two parts: a fitting set, and a held out set. The recommenders are then trained on the fitting set. Then, the recommenders predict the labels of the

held out set. The optimal alpha can either minimize the mean absolute error

$$\mathcal{E}_1 = \frac{1}{m} \left\| R - \vec{R}\vec{\alpha} \right\|_1$$

or the (half) mean squared error

$$\mathcal{E}_2 = \frac{1}{2m} \left\| R - \vec{R}\vec{\alpha} \right\|_2^2$$

m as the number of samples in the held out set, $R = (r_1, \dots, r_m)^\top$ as the vector of true ratings of the held out samples, $\vec{R} = (\vec{r}_1, \dots, \vec{r}_m)^\top$ as the matrix of predictions on the held out samples.

$\vec{\alpha}$ is then optimized with gradient descent. Every entry is initialized with $\frac{1}{K}$. At each step, $\vec{\alpha} \leftarrow \vec{\alpha} - \gamma \nabla \mathcal{E}(\vec{\alpha})$ is pushed into the direction of the negative gradient of the error, scaled with a learning rate γ . This is done until the error changes less than a threshold ϵ , or the maximum amount of iterations is reached.

$$\begin{aligned} \nabla \mathcal{E}_1(\vec{\alpha}) &= \frac{1}{m} \text{sign} \left(R - \vec{R}\vec{\alpha} \right)^\top \vec{R} \\ \nabla \mathcal{E}_2(\vec{\alpha}) &= \frac{1}{m} \left(R - \vec{R}\vec{\alpha} \right)^\top \vec{R} \end{aligned}$$

After $\vec{\alpha}$ has been determined, the predictors are trained on the whole training set to not miss out on information. Special cases do not need to be handled, as they are handled in the individual recommenders.

2 Results

The hyper parameters were tuned based on the root mean squared error and the mean absolute error from cross validation. Those were the best performing ones:

- user based:
 - $k = 50$ neighbours in peer group
 - $\alpha = 1$, not amplifying importance of similarity
 - $S_{min} = 0.4$ minimum similarity
 - $\beta = 4$ common dimensions before similarity reduction
 - 8 clusters
- item based:
 - $k = 50$ neighbours in peer group
 - $\alpha = 1$, not amplifying importance of similarity
 - $S_{min} = 0.6$ minimum similarity
 - $\beta = 6$ common dimensions before similarity reduction
- factorization based:
 - $\lambda = 5$ regularization parameter
 - $f = 20$ latent dimensions
 - biases not fixed
 - $\epsilon = 10^{-2}$, max 20 iterations
- hybrid of user based, item based, factorization based:
 - $\gamma = 1$ learning rate
 - $\epsilon = 10^{-4}$, max 100 iterations




Figure 1: Hybrid error during gradient descent

It is worthy to note that the hybrid didn't improve much compared to the initial equal weights. In [Figure 1](#) you can see mean absolute error of the hybrid recommender during the gradient descent of $\vec{\alpha}$.

2.1 Performance

For time measurement, the recommenders were trained on `train.csv` (435K samples) and evaluated on `qualifying_blanc.csv` (109K samples).

Recommender	offline phase	online phase
user based	15.97s	6.84s
item based	44.3s	4.79s
factorization	85.65s	0.53s
hybrid	304.1s	26.79s

Table 1: Runtime in offline and online phase

2.2 Error Scores

The error scores were measured with cross validation on the training set, with $R = 8$ rotations. Root mean squared error (rmse) and mean absolute error (mae) were calculated.

Recommender	rmse	mae
user based	0.670	0.301
item based	0.569	0.222
factorization	0.559	0.233
hybrid	0.513	0.202

Table 2: Error scores

2.3 Final Score

My final score on the `validation set` is 0.530. It was achieved with the hybrid recommender.